
Gradient Descent on Linear Regression

A Quick Summary of Stochastic, Batch and Mini-batch Gradient Descent

Hair Albeiro Parra Barrera
March 5, 2020

1 Gradient Descent for Linear Regression

Suppose we have a hypothesis $h : \mathbb{R}^n \rightarrow \mathbb{R}$, $h_\theta(\mathbf{x}) = \hat{y}$ with parameters $\theta \in \mathbb{R}^n$. Recall the **Mean-Squared Loss (MSE)** metric, applied to linear regression:

$$MSE(y, h_\theta(\mathbf{x})) = \frac{1}{2n} \|\mathbf{y} - h_\theta(\mathbf{x})\|_2^2 = \frac{1}{2n} \sum_{i=1}^n (y^{(i)} - \hat{y}^{(i)})^2 = \frac{1}{2n} \sum_{i=1}^n (y^{(i)} - \mathbf{w}^T \mathbf{x}^{(i)})^2$$

In order to minimize the MSE, we take partials w.r.t. each parameter, and have the general update:

$$\mathbf{w}^{k+1} = \mathbf{w}^k + \alpha \frac{\partial}{\partial \mathbf{w}} MSE(y, \hat{y})$$

As the MSE is a **convex function**, it is guaranteed to have a global optimum, so given an appropriate choice of α , also called the **learning rate**, the algorithm will converge. In what follows, the algorithm stops whenever $\|\mathbf{w}^k - \mathbf{w}^{k-1}\| < \epsilon$ or $|MSE_k - MSE_{k-1}| < \epsilon$, for some small ϵ .

Batch Gradient Descent: For $k = 0, 1, \dots$

1. For $k = 0, 1, \dots$

$$\mathbf{w}^{k+1} = \mathbf{w}^k + \alpha \frac{1}{n} \sum_{i=1}^n (y^{(i)} - \mathbf{w}^T \mathbf{x}^{(i)}) \mathbf{x}^{(i)}$$

Mini-batch Gradient Descent:

1. For $k = 0, 1, \dots$

- (a) Split data D into T subsets D_t of sizes n_0, \dots, n_{T-1} , s.t. $\sum_t n_t = 1$.
- (b) For each subset D_t :

$$\mathbf{w} := \mathbf{w} + \alpha \frac{1}{n_t} \sum_{i: \mathbf{x}^{(i)} \in D_t}^{n_t} (y^{(i)} - \mathbf{w}^T \mathbf{x}^{(i)}) \mathbf{x}^{(i)}$$

Stochastic Gradient Descent:

1. For $k = 0, 1, \dots$

- (a) For $i = 1, \dots, n$:

$$\mathbf{w} := \mathbf{w} + \alpha (y^{(i)} - \mathbf{w}^T \mathbf{x}^{(i)}) \mathbf{x}^{(i)}$$